

Sorted Ordered Column-family Stores

Google's Bigtable, HBase, Hypertable

Shashank Tiwari

blog: shanky.org | twitter: @tshanky

st@treasuryofideas.com

TreasuryOfIdeas

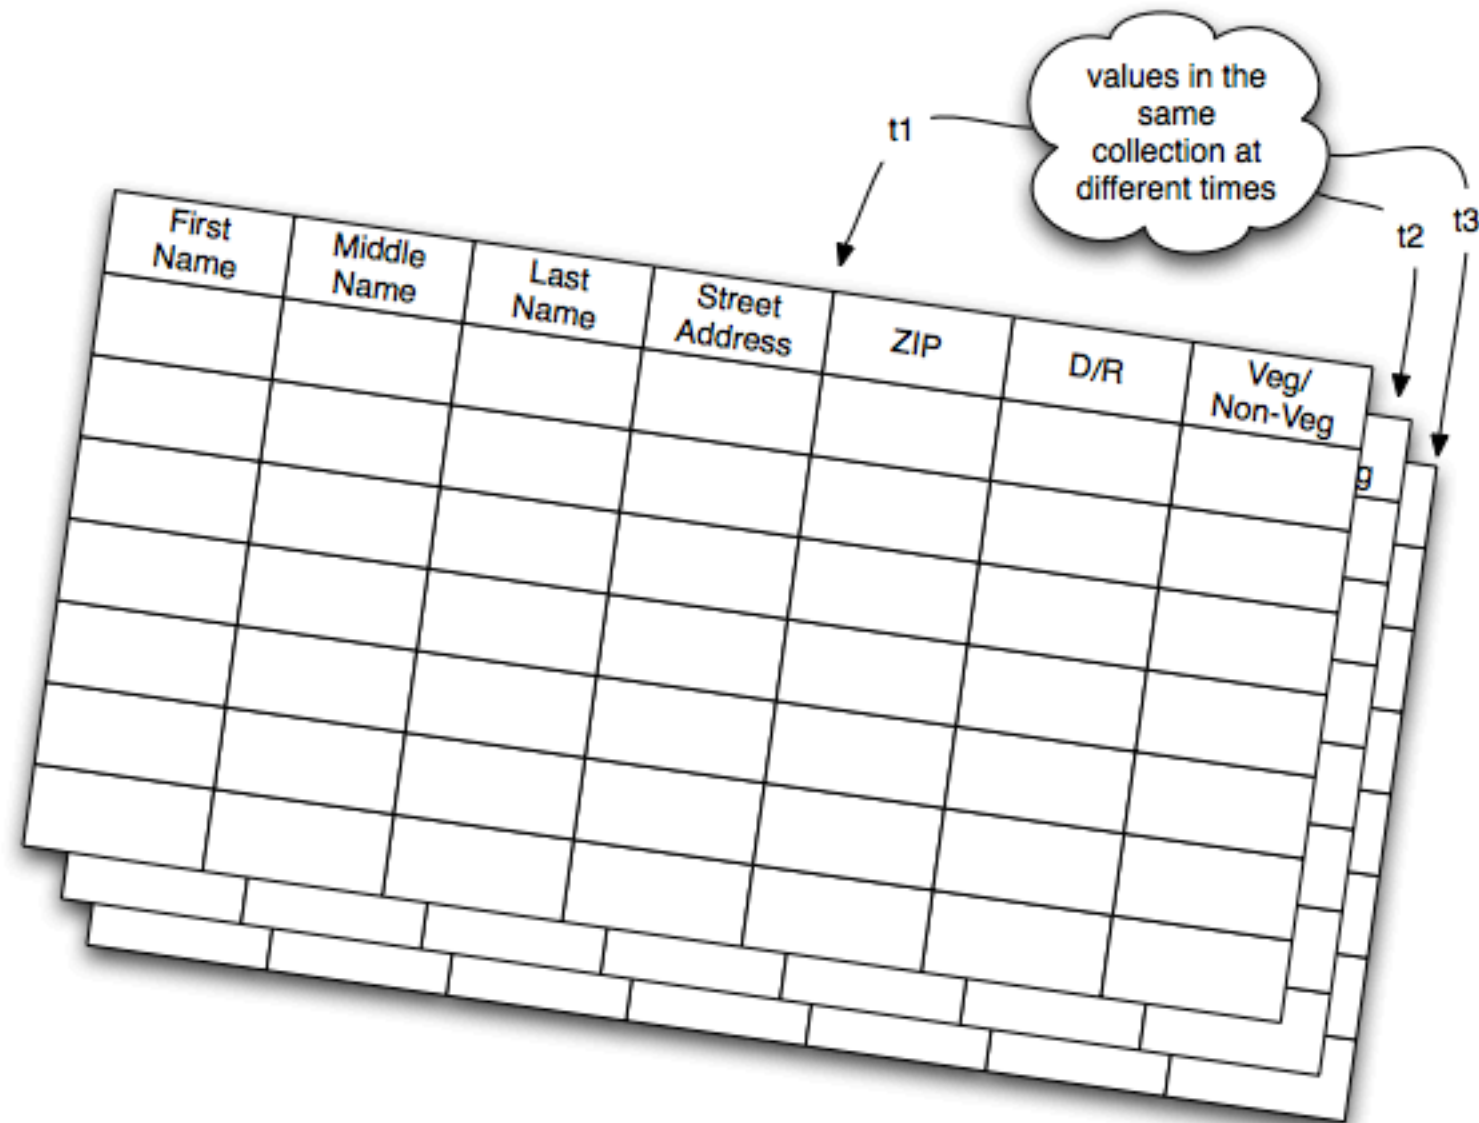
Definition

- Sorted
- Ordered
- Distributed
- Map

Essential schema

| row key | time | contents | anchor | mime |
|-------------|------|-------------|------------|-------------|
| com.cnn.www | t9 | | cnnsi.com | |
| | t8 | | my.look.ca | |
| | t6 | "<html>..." | | "text/html" |
| | t5 | "<html>..." | | |
| | t3 | "<html>..." | | |
| | | | | |

Multi-dimensional View



A Map/Hash View

```
{
```

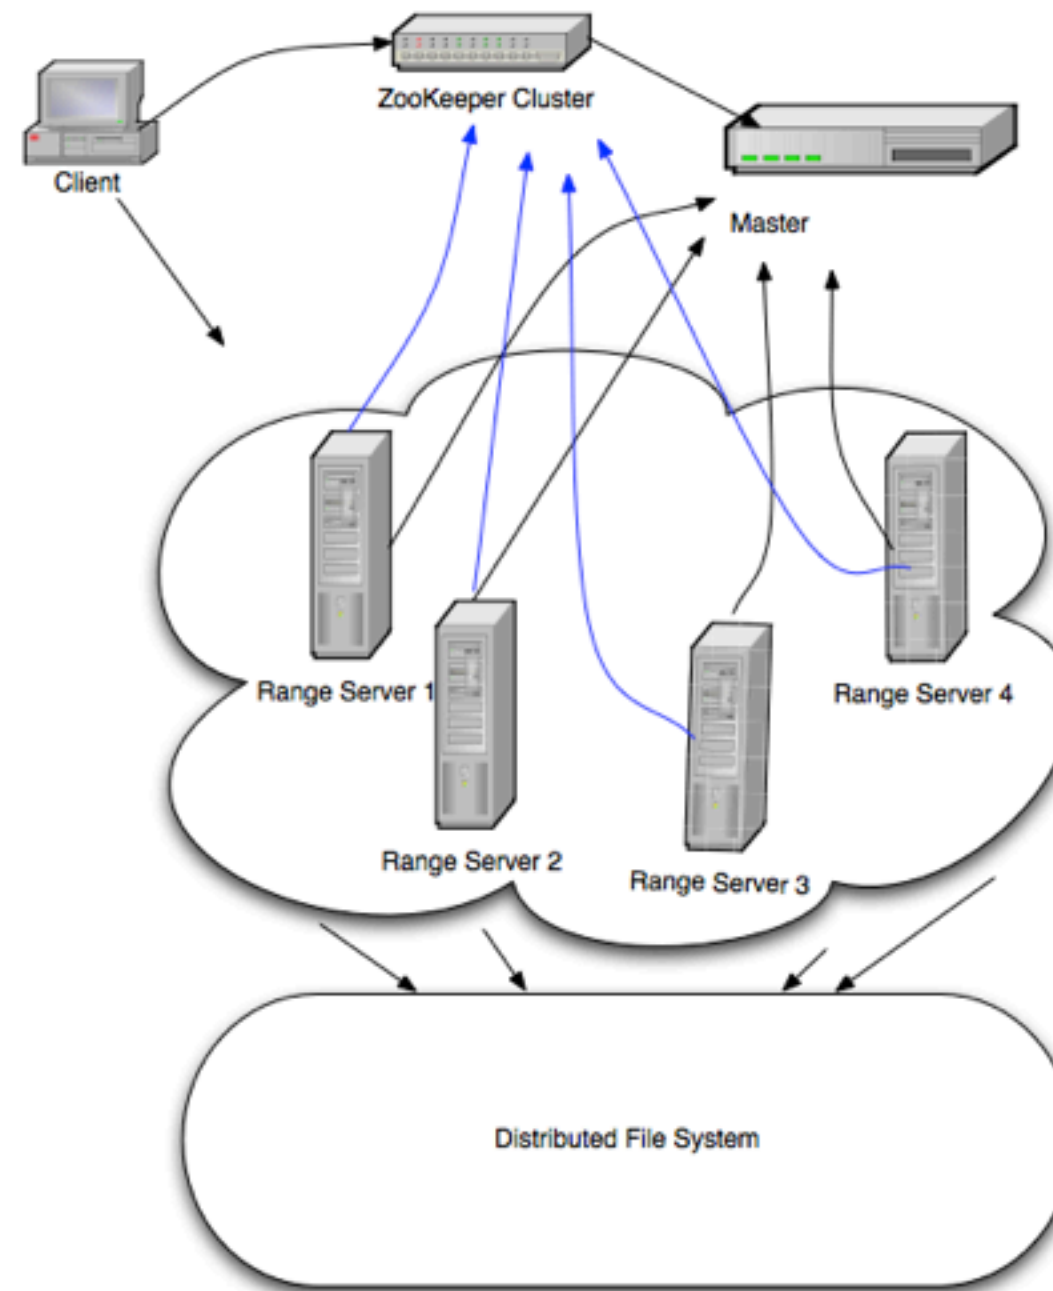
```
  "row_key_1" : { "name" : {
```

```
    "first_name" : "Jolly", "last_name" : "Goodfellow"
```

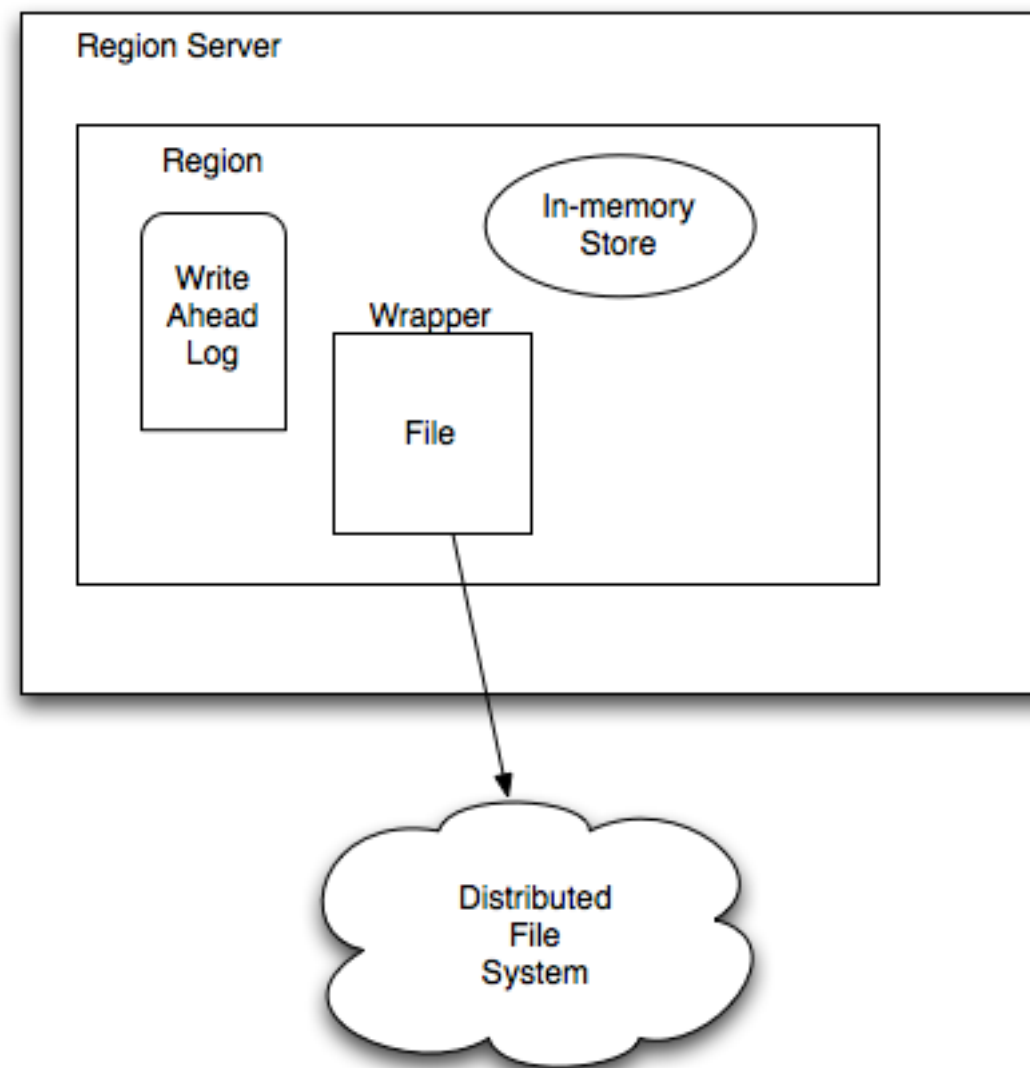
```
  } } },
```

```
  "location" : { "zip": "94301" },
```

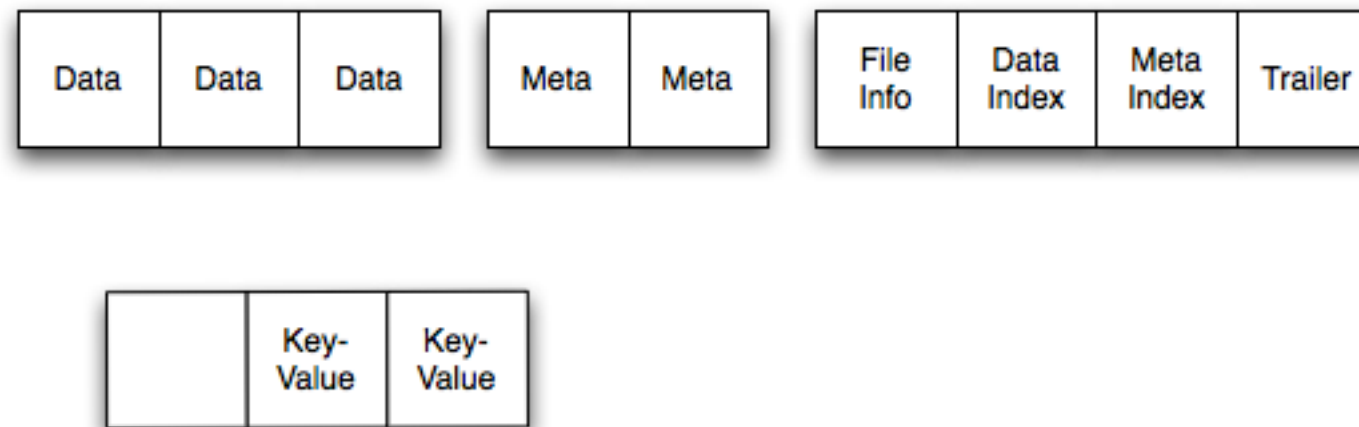
Architectural View (HBase)



The Persistence Mechanism



The underlying file format



Model Wrappers (The GAE Way)

- Python
 - Model, Expando, PolyModel
- Java
 - JDO, JPA

HBase Data Access

- Thrift + Avro
- Java API -- HTable, HBaseAdmin
- Hive (SQL like)
- MapReduce -- sink and/or source

Transactions

- Atomic row level
- GAE Entity Groups

Indexes

- Row ordered
- Secondary indexes
- GAE style multiple indexes
 - thinking from output to query

Use cases

- Many Google's Products
- Facebook Messaging
- StumbleUpon
- Open TSDB
- Mahalo, Ning, Meetup, Twitter, Yahoo!
- Lily -- open source CMS built on HBase & Solr

Questions?

blog: shanky.org | twitter: @tshanky

st@treasuryofideas.com

TreasuryOfIdeas